

BookKeeper Getting Started Guide

by

Table of contents

| | |
|--------------------------------------|---|
| 1 Programming with BookKeeper..... | 2 |
| 1.1 Instantiating BookKeeper..... | 2 |
| 1.2 Creating a ledger. | 2 |
| 1.3 Adding entries to a ledger. | 3 |
| 1.4 Closing a ledger. | 4 |
| 1.5 Opening a ledger. | 5 |
| 1.6 Reading from ledger | 6 |
| 1.7 Deleting a ledger | 6 |

1 Programming with BookKeeper

- [Instantiating BookKeeper.](#)
- [Creating a ledger.](#)
- [Adding entries to a ledger.](#)
- [Closing a ledger.](#)
- [Opening a ledger.](#)
- [Reading from ledger](#)
- [Deleting a ledger](#)

1.1 Instantiating BookKeeper.

The first step to use BookKeeper is to instantiate a BookKeeper object:

```
org.apache.bookkeeper.BookKeeper
```

There are three BookKeeper constructors:

```
public BookKeeper(String servers) throws KeeperException,  
IOException
```

where:

- `servers` is a comma-separated list of ZooKeeper servers.

```
public BookKeeper(ZooKeeper zk) throws InterruptedException,  
KeeperException
```

where:

- `zk` is a ZooKeeper object. This constructor is useful when the application also using ZooKeeper and wants to have a single instance of ZooKeeper.

```
public BookKeeper(ZooKeeper zk, ClientSocketChannelFactory  
channelFactory) throws InterruptedException, KeeperException
```

where:

- `zk` is a ZooKeeper object. This constructor is useful when the application also using ZooKeeper and wants to have a single instance of ZooKeeper.
- `channelFactory` is a netty channel object
(`org.jboss.netty.channel.socket`).

1.2 Creating a ledger.

Before writing entries to BookKeeper, it is necessary to create a ledger. With the current BookKeeper API, it is possible to create a ledger both synchronously or asynchronously. The following methods belong to `org.apache.bookkeeper.client.BookKeeper`.

Synchronous call:

```
public LedgerHandle createLedger(int ensSize, int qSize,
DigestType type, byte passwd[]) throws KeeperException,
InterruptedException, IOException, BKException
```

where:

- `ensSize` is the number of bookies (ensemble size);
- `qSize` is the write quorum size;
- `type` is the type of digest used with entries: either MAC or CRC32.
- `passwd` is a password that authorizes the client to write to the ledger being created.

All further operations on a ledger are invoked through the `LedgerHandle` object returned.

As a convenience, we provide a `createLedger` with default parameters (3,2,VERIFIABLE), and the only two input parameters it requires are a digest type and a password.

Asynchronous call:

```
public void asyncCreateLedger(int ensSize, int qSize,
DigestType type, byte passwd[], CreateCallback cb, Object
ctx )
```

The parameters are the same of the synchronous version, with the exception of `cb` and `ctx`. `CreateCallback` is an interface in `org.apache.bookkeeper.client.AsyncCallback`, and a class implementing it has to implement a method called `createComplete` that has the following signature:

```
void createComplete(int rc, LedgerHandle lh, Object ctx);
```

where:

- `rc` is a return code (please refer to `org.apache.bookkeeper.client.BKException` for a list);
- `lh` is a `LedgerHandle` object to manipulate a ledger;
- `ctx` is a control object for accountability purposes. It can be essentially any object the application is happy with.

The `ctx` object passed as a parameter to the call to create a ledger is the one same returned in the callback.

1.3 Adding entries to a ledger.

Once we have a ledger handle `lh` obtained through a call to create a ledger, we can start writing entries. As with creating ledgers, we can write both synchronously and asynchronously. The following methods belong to `org.apache.bookkeeper.client.LedgerHandle`.

Synchronous call:

`public long addEntry(byte[] data) throws InterruptedException`
 where:

- `data` is a byte array;

A call to `addEntry` returns the status of the operation (please refer to `org.apache.bookkeeper.client.BKDefs` for a list);

Asynchronous call:

`public void asyncAddEntry(byte[] data, AddCallback cb, Object ctx)`

It also takes a byte array as the sequence of bytes to be stored as an entry. Additionally, it takes a callback object `cb` and a control object `ctx`. The callback object must implement the `AddCallback` interface in `org.apache.bookkeeper.client.AsyncCallback`, and a class implementing it has to implement a method called `addComplete` that has the following signature:

`void addComplete(int rc, LedgerHandle lh, long entryId, Object ctx);`

where:

- `rc` is a return code (please refer to `org.apache.bookkeeper.client.BKDefs` for a list);
- `lh` is a `LedgerHandle` object to manipulate a ledger;
- `entryId` is the identifier of entry associated with this request;
- `ctx` is control object used for accountability purposes. It can be any object the application is happy with.

1.4 Closing a ledger.

Once a client is done writing, it closes the ledger. The following methods belong to `org.apache.bookkeeper.client.LedgerHandle`.

Synchronous close:

`public void close() throws InterruptedException`

It takes no input parameters.

Asynchronous close:

`public void asyncClose(CloseCallback cb, Object ctx) throws InterruptedException`

It takes a callback object `cb` and a control object `ctx`. The callback object must implement the `CloseCallback` interface in

`org.apache.bookkeeper.client.AsyncCallback`, and a class implementing it has to implement a method called `closeComplete` that has the following signature:

```
void closeComplete(int rc, LedgerHandle lh, Object ctx)
```

where:

- `rc` is a return code (please refer to `org.apache.bookkeeper.client.BKDefs` for a list);
- `lh` is a `LedgerHandle` object to manipulate a ledger;
- `ctx` is control object used for accountability purposes.

1.5 Opening a ledger.

To read from a ledger, a client must open it first. The following methods belong to `org.apache.bookkeeper.client.BookKeeper`.

Synchronous open:

```
public LedgerHandle openLedger(long lId, DigestType type, byte
passwd[]) throws InterruptedException, BKException
```

- `ledgerId` is the ledger identifier;
- `type` is the type of digest used with entries: either `MAC` or `CRC32`.
- `passwd` is a password to access the ledger (used only in the case of `VERIFIABLE` ledgers);

Asynchronous open:

```
public void asyncOpenLedger(long lId, DigestType type, byte
passwd[], OpenCallback cb, Object ctx)
```

It also takes a ledger identifier and a password. Additionally, it takes a callback object `cb` and a control object `ctx`. The callback object must implement the `OpenCallback` interface in `org.apache.bookkeeper.client.AsyncCallback`, and a class implementing it has to implement a method called `openComplete` that has the following signature:

```
public void openComplete(int rc, LedgerHandle lh, Object ctx)
```

where:

- `rc` is a return code (please refer to `org.apache.bookkeeper.client.BKDefs` for a list);
- `lh` is a `LedgerHandle` object to manipulate a ledger;
- `ctx` is control object used for accountability purposes.

1.6 Reading from ledger

Read calls may request one or more consecutive entries. The following methods belong to `org.apache.bookkeeper.client.LedgerHandle`.

Synchronous read:

```
public Enumeration<LedgerEntry> readEntries(long firstEntry,
long lastEntry) throws InterruptedException, BKException
```

- `firstEntry` is the identifier of the first entry in the sequence of entries to read;
- `lastEntry` is the identifier of the last entry in the sequence of entries to read.

Asynchronous read:

```
public void asyncReadEntries(long firstEntry, long
lastEntry, ReadCallback cb, Object ctx) throws BKException,
InterruptedException
```

It also takes a first and a last entry identifiers. Additionally, it takes a callback object `cb` and a control object `ctx`. The callback object must implement the `ReadCallback` interface in `org.apache.bookkeeper.client.AsyncCallback`, and a class implementing it has to implement a method called `readComplete` that has the following signature:

```
void readComplete(int rc, LedgerHandle lh,
Enumeration<LedgerEntry> seq, Object ctx)
```

where:

- `rc` is a return code (please refer to `org.apache.bookkeeper.client.BKDefs` for a list);
- `lh` is a `LedgerHandle` object to manipulate a ledger;
- `seq` is a `Enumeration<LedgerEntry>` object to containing the list of entries requested;
- `ctx` is control object used for accountability purposes.

1.7 Deleting a ledger

Once a client is done with a ledger and is sure that nobody will ever need to read from it again, they can delete the ledger. The following methods belong to `org.apache.bookkeeper.client.BookKeeper`.

Synchronous delete:

```
public void deleteLedger(long lId) throws
InterruptedException, BKException
```

- `lId` is the ledger identifier;

Asynchronous delete:

```
public void asyncDeleteLedger(long lId, DeleteCallback cb,  
Object ctx)
```

It takes a ledger identifier. Additionally, it takes a callback object `cb` and a control object `ctx`. The callback object must implement the `DeleteCallback` interface in `org.apache.bookkeeper.client.AsyncCallback`, and a class implementing it has to implement a method called `deleteComplete` that has the following signature:

```
void deleteComplete(int rc, Object ctx)
```

where:

- `rc` is a return code (please refer to `org.apache.bookkeeper.client.BKDefs` for a list);
- `ctx` is control object used for accountability purposes.